Deep learning.

- Model parametrization
- Defining a loss function
- Optimization
- Model selection
- Reporting

Parametrization

$p(y|x)$

output input

① $p(y|x) \propto \exp\{F(x,y)\}$

② latent variable approach

$p(y,z|x) \propto \exp\{F(x,y,z)\}$

$p(z|x) \propto \exp\{F'(x,z)\}$

Multi class classification.

$y \in \{1, \cdots, C\}$

$F(x,y) = \underset{\uparrow \mathbb{R}^d}{w_y^T} \underset{\uparrow \mathbb{R}^d}{\phi(x)} + \underset{\uparrow \mathbb{R}}{b_y}$

$$p(y|x) = \frac{\exp\{w_y^T \phi(x) + b_y\}}{\sum_{y'=1}^{C} \exp\{w_{y'}^T \phi(x) + b_y\}}$$

Mixture density network

$y \in \mathbb{R}^{d'}$

$z \in \{1, \cdots, C\}$

$W_\mu^z \in \mathbb{R}^{d' \times d}$, $\phi(x) \in \mathbb{R}^d$

$$F(x, y, z) = -\| y - W_\mu^z \phi(x) \|^2$$

$$F(x, z) = W_z^T \phi(x) + b_z, \quad W_z \in \mathbb{R}^d, \; b_z \in \mathbb{R}$$

$$p(y|x) = \sum_{z=1}^{C} p(y, z|x) = \sum_{z=1}^{C} p(y|z, x) p(z|x)$$

$$= \sum_{z=1}^{C} \left(\frac{\exp\{W_z^T \phi(x) + b_z\}}{\sum_{z'=1}^{C} \exp\{W_{z'}^T \phi(x) + b_{z'}\}} \right) \cdot \left(\frac{\exp\{-\| y - W_\mu^z \phi(x) \|^2\}}{\int_{\mathbb{R}^{d'}} \exp\{-\| y' - W_\mu^z \phi(x) \|^2\} dy'} \right)$$

Loss function

(per-example loss)

for (x, y^*), l is $-\log p(y^*|x)$

for $D = \{(x_1, y_1^*), \cdots, (x_N, y_N^*)\}$,

$$p(D) = \prod_{n=1}^{N} p(x_n, y_n^*) = \prod_{n=1}^{N} p(y_n^*|x_n) \underbrace{p(x_n)}_{\text{uniform}}$$

$$\log p(D) = \sum_{n=1}^{N} \log p(y_n^*|x_n) + \text{const.}$$

Multiway classification

$$-\log p(y^*|x) = \underbrace{-W_{y^*}^T \phi(x) - b_{y^*}}_{= F(x, y^*)} + \log \sum_{y'=1}^{C} \exp\{W_{y'}^T \phi(x) + b_{y'}\}$$

Optimization : Gradient descent

$L(D;\theta)$: a data-level loss func.

$$L(D;\theta) = \frac{1}{|D|} \sum_{(x,y)\in D} l(x,y;\theta)$$

$$\theta \leftarrow \theta - \eta \underbrace{\nabla_\theta L(D;\theta)}_{\text{gradient}} = \theta - \frac{\eta}{|D|} \underbrace{\sum_{x,y\in D} \nabla_\theta l(x,y;\theta)}$$

for $(x,y)\in D$, $x,y \sim P_D$ $\quad M \subset D$ (minibatch) $\quad |M| \ll |D|$

$$\nabla_\theta L(D;\theta) \simeq \mathbb{E}_{x,y\sim P_D}[\nabla l(x,y;\theta)] \simeq \frac{1}{|M|} \sum_{(x,y)\in M} \nabla l(x,y;\theta)$$

Stochastic Gradient Descent

Model selection

= hyperparameter selection/optimization

$\lambda \in \Lambda$: a hyperparameter

$$\min_{\lambda\in\Lambda} \mathbb{E}_{(x,y)\sim D'} \left[l'\left(x,y ; \operatorname{argmin}_\theta \mathbb{E}_{(x,y)\sim D}[l(x,y;\theta)]\right)\right]$$

validation set D' ; training set D

Reporting