a language model that can handle unbounded context (prefix).

$$p(x_1, x_2, \cdots, x_T) = \prod_{t=1}^{T} p(x_t \mid x_1, \cdots, x_{t-1})$$

$x_t \in V$

$\approx p(x_t \mid x_{t-n}, \cdots, x_{t-1})$: n-gram LM

Neural probabilistic language model.

① $e(x_{t-n}, \cdots, x_{t-1}) = [e(x_{t-n})^T, \cdots, e(x_{t-1})^T]^T \in \mathbb{R}^{nd}$

② $h = \tanh(W e(x_{t-n}, \cdots, x_{t-1}) + b) \in \mathbb{R}^{d'}$

③ $F(x_t, x_{t-n:t-1}) = e'(x_t)^T h + b'(x_t)$

$W \in \mathbb{R}^{d' \times nd}$

$W \in \mathbb{R}^{d' \times n1}$ class embedding

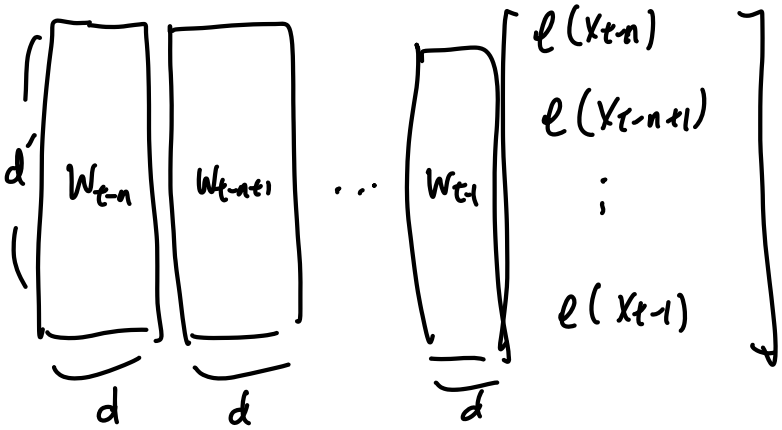

$$= \sum_{n'=1}^{n} W_{t-n'} e(x_{t-n'})$$

A text classifier summarizes the input text

$X = (x_1, \cdots, x_T)$

$\phi(X) = \sigma(W \sum_{t=1}^{T} e(x_t) + b)$, $F(X, y) = w_y^T \phi(X) + c_y$

A recurrent neural network.

: recursively compress the input sequence.

If $h_{t-1} = \phi(x = [w_1, \dots, w_{t-1}])$, can we add w_t?

- $h_t = R(h_{t-1}, w_t)$, where e.g. $R(h_{t-1}, w_t)$

$\uparrow \mathbb{R}^{d'}$ $\uparrow \mathbb{R}^{d'}$

$= \sigma\left(W_r h_{t-1} + W_x e(w_t) + b \right)$

$\uparrow \mathbb{R}^{d' \times d'}$ $\uparrow \mathbb{R}^{d'}$ $\uparrow \mathbb{R}^{d' \times d}$ $\uparrow \mathbb{R}^{d}$ $\uparrow \mathbb{R}^{d'}$

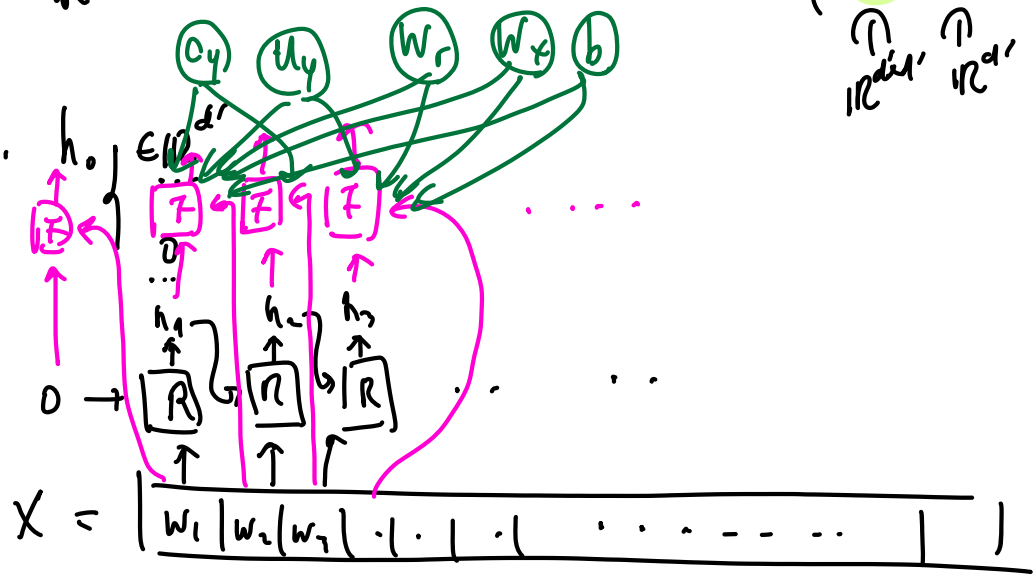

A recurrent neural net language model [Mikolov et al. 2010]

Alice $\Longrightarrow$ Bob

$m \in \{1, \dots, c\}$

$\log_2 c$ bits

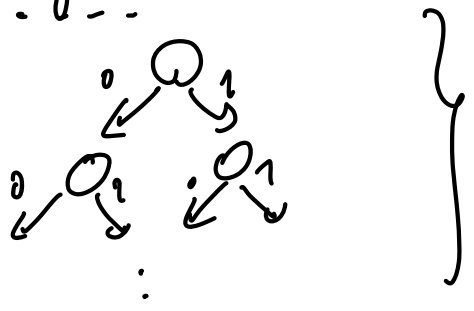

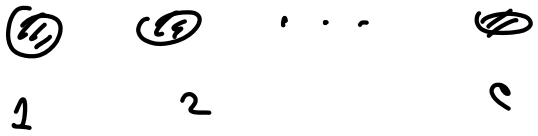

$- \sum p(m) \log_2 p(m)$

$2^{-\log_2 p(m)}$ ⇒ $2^{-\log_2 p(w_t | w_{<t})}$

$$B = -\frac{1}{\sum_{n=1}^{N} T^n} \sum_{n=1}^{N} \sum_{t=1}^{T^n} \log_2 p(w_t^n | w_{<t}^n) \quad , \text{ perplexity} = 2^B$$